

Glomerulus Segmentation in Kidney Biopsies

Computer Vision Project, Group 9

Lukas Göbl, Peer Schäfer, Lukas Scheib


15.07.2025

Problem Statement


General: Accurate segmentation of glomeruli in kidney biopsy images is crucial for diagnosing chronic kidney diseases.

Precise Tasks: Segment glomeruli at patch-level and whole-slide level in kidney tissue images as part of the MICCAI 2024 KPIs Challenge.

WSI Example: train\56Nx\12_116



WSI Mask Example: train\56Nx\12_116



- Standard UNet architecture with double convolution layers and skip connections.
- Dilated convolution in the bottleneck (with dilations 1, 2 and 4 with $\frac{1}{3}$ of the channels each) for multi-scale feature extraction.¹
- Channel counts scale with variable patch size.

¹Li et al. (2021)

- Trained on official KPIs dataset (train/val/test splits).
- Combined BCE and Dice loss (BCEDiceLoss) to address class imbalance.
- Hyperparameter tuning: BCE/Dice ratio, Dice smooth, positive class weighting for BCE, base learning rate of Adam optimizer.
- Experiments with different patch sizes (256, 2048 px) and data subsets.
- Evaluated using accuracy, precision, recall, and F_1 -score.
- Qualitative assessment via visual inspection.

Metrics of our best performing Models

Patch size	Split	Accuracy	Precision	Recall	F_1 -score
256	Validation	0.9781	0.7479	0.7660	0.7569
	Test	0.9757	0.7537	0.7133	0.7330
2048	Validation	0.9705	0.9299	0.3662	0.5255
	Test	0.9643	0.9159	0.2616	0.4070

Hyperparameters:

- BCEDiceLoss Ratio = 0.8
- BCE Positive Weight = 25
- Dice Smooth = 10^{-6}
- Learning Rate = 10^{-4}

Implemented patching pipeline with the following hyperparameters:

- Patch size
- Patch overlap (stride)
- Trivial patch threshold for positive pixels

Possible next steps:

- Implementation of a stitching pipeline to combine patch predictions into a full WSI mask.
- Train VariableUNet using the complete patching and stitching pipeline.

What worked:

- BCE + Dice loss crucial for recall and F_1 -score.
- Hyperparameter tuning improved glom detection.
- WSIs can be patched for patch-level segmentation.

Limitations:

- Instability in training, especially with full-size patches.
- Our configurations, trained on small patches, did not generalize well to larger patches.
- No empirical validation on WSI segmentation due to time constraints.

In retrospect:

- Greater emphasis on pre- and post-processing (downsampling patches and upsampling predictions).
- More in-depth data exploration regarding samples of different diseases.

Future work:

- Exploring different model architectures and methods.
- Search for more robust hyperparameters and domain adaptation.
- Continued work on WSI segmentation pipeline.

Appendix

UNetArchitecture





Figure: Doi: 10.1117/1.JMI.8.6.067501


Double Convolution

```
1 class DoubleConv(nn.Module):
2     def __init__(self, in_channels, out_channels):
3         super().__init__()
4         self.double_conv = nn.Sequential(
5             nn.Conv2d(in_channels, out_channels, 3,
6                       padding=1),
7             nn.BatchNorm2d(out_channels),
8             nn.ReLU(inplace=True),
9             nn.Conv2d(out_channels, out_channels, 3,
10                      padding=1),
11             nn.BatchNorm2d(out_channels),
12             nn.ReLU(inplace=True),
13         )
14
15     def forward(self, x):
16         return self.double_conv(x)
```

Heat Map F1 Score



Heat Map Recall



FlipFlop Graph

